





# Merged proofs in Verkle tree and optimization




Ressac\_No1

# Acknowledgments

- My appreciation of their instruction and discussion with me:
  - EF Stateless consensus & Verkle-tree migration team
    -  Guillaume Ballet
    -  Ignacio Hagopian
    -  Josh Rudolf
  -  g11tech
  - Milos Stankovic
  - Dankrad Feist

# From Merkle to Verkle


- Merkle-Patricia Trie has overhead in proof size



- In a MPT of width  $d$  and branch width  $2^q$
- The proof of the value of key  $k[0:d]$  consists of  $d$  hashing computations, each convinces that a node along the root-leaf path has the correct hash encoded from its all children
- Thus, the hashes of all the nodes along the path and their siblings (except the root) must be contained in the proof
- Total proof size  $O(2^q d)$

# From Merkle to Verkle

- Verkle tree reduces the proof size by avoiding sibling hashes



- A node's hash is derived from the commitment of the polynomial with children's hashes as evaluations at point  $0, 1, \dots, 2^q-1$
- The proof only contains those hashes of nodes along the root-leaf path and correctness of every evaluation:  $H[C_{i+1}, k[i]] = f_{i,k[i-1]}(k[i])$  in which  $f_{i,k[i-1]}$  is committed as  $C_{i,k[i-1]}$


# Verkle proof of one node

- Given a key  $k[0:d]$ , along its access path in the Verkle tree, it must prove for each node at depth  $i \in [0:d-1]$  that  $H[C_{i+1}, k[i]] = f_{i,k[i-1]}(k[i])$ , in which  $f_{i,k[i-1]}$  is committed as  $C_{i,k[i-1]}$
- Polynomial commitment scheme: ~~KZG ( $O(1)$  proof size,  $O(1)$  pairing check in verification)~~ Inner Product Arguments ( $O(\log n)$  proof size,  $O(\log n)$  group operations in verification)
- KZG is discarded because of the compulsory trusted setup
- Pedersen+IPA,  $n$  is the degree of the committed polynomial, equal to branch width ( $2^q$ ) in Verkle tree proof.

# Pedersen+IPA multiproof


- Pedersen commitment is additive homomorphic
  - An  $n$ -dimension vector  $x$  is committed as  $C_x = \sum_{i=0}^{n-1} x_i g_i$
  - In which  $g$  is an  $n$ -dimension vector of group elements, as the basis
- Hence, the opening of multiple commitments can be merged
  - To show  $f_0(z_0) = y_0, f_1(z_1) = y_1, \dots, f_{d-1}(z_{d-1}) = y_{d-1}$
  - Enough to show  $\sum_{j=0}^{d-1} a_j f_j(z_j) = \sum_{j=0}^{d-1} a_j y_j$  in case all  $a_j$  are random
  - Commit the linear combined  $2^q d$ -dimension vector and open it via IPA
  - $O(q \log d)$  proof size and verification time for a SLOAD access

# Merged proof: shared commitments



- In multiple accesses of the same leaf node, naturally, the proof remains the same
- In accesses of multiple nodes? Some ancestors are shared
- Example: two keys  $k$  and  $k'$  differs only at the last letter

# Merged proof: shared commitments



One proof:  $d$  commitments involved

Two proofs:  $d+1$  commitments involved

- In multiple accesses of the same leaf node, naturally, the proof remains the same
- In accesses of multiple nodes? Some ancestors are shared
- Example: two keys  $k$  and  $k'$  differs only at the last letter



# Stem and suffix

- In the Verkle tree scheme applied to Ethereum upgrade,  $q=8$ , each edge matches a byte
- A key's last byte is its **suffix** and the left is its **stem**
- Basic form of proofs of storage slot operations is up to stems, and the commitments of sibling leaf nodes (representing the same stem but different suffices) are opened together
- Amortized gas cost is much lower to access multiple keys with shared stems (EIP-4762)


# Merged proof: rebuilding the witness tree

- In practice, it is the block builder to determine the accessed (including updated) storage slots caused by the txns in the block
- Clients are stateless, only trusted block builders own the world states, including codes and stored values in the Verkle tree
- Rebuilding the “witness tree” of all accessed nodes and their values and commitments, with updates as the result of SSTOREs
- Then producing proofs of EVM execution transcripts and storage operations based on the Verkle tree storage

# But how to deal with SSTORE?


- A SSTORE in a block modifies the polynomial and commitment of every node along the root-leaf path
- Because the root commitment is also updated, old committed values in the “partial witness tree” seem to expire after a SSTORE
- But, really?

# Part of old multiproofs still available



Partial witness  
tree before  
SSTORE


# Part of old multiproofs still available



Partial witness  
tree **after**  
SSTORE

- In the multiproof of the green leaf node, all the ancestors other than the root remains the same committed polynomials after a SSTORE at the red leaf node
- This part is still available for proofs after the SSTORE

# Slices of multiproof



- Consider a SSTORE at the orange leaf followed by a SSTORE at the red leaf
- The proof of the green leaf can be described as three slices (root, orange one at depth 1, the other nodes)
- A slice is linearly combined into a multiproof and reusable

# Optimized redesigning: proofs of SSTORE?

- It seems ultimate to optimize a Verkle proof that makes use of multiproofs and slices
- Bottleneck: a Verkle proof of SSTORE must involve  $O(d)$  polynomial commitments, indicating all the nodes along a root-leaf path
- **Our goal: to redesign the Eth storage system to reach both client statelessness and  $o(d)$  proof size and verification time for a SSTORE (may be amortized)**
- Idea: to record events and proofs of SSTORE?

# Witness of SSTORE

- The block builder can organize all the SSTORE (and other storage update) events in a witness vector/tree/other ADS
- An element of the witness of SSTORE contains the targeted storage slot key and value, as well as its timestamp
- To prove a SLOAD
  - If updated by a SSTORE in the witness, prove its value correctness and its timestamp before this SLOAD with no other SSTORE at the same slot in between
  - If not updated by any SSTORE in the witness, produce a proof from the world state



# Frequent updated slots

- Observation: a few storage slots are frequently updated, within one block, or even within one txn
- E.g. counters, nonces and account balances
- Should be experiments to confirm this observation
- Shorter proofs of SSTORE and SLOAD events in building a block, and update in the world state only once per slot per block

# Optional: storage slots with small values

- This optimization idea was proposed in [a post by Milos](#)
- MPT and the adopted Verkle tree scheme distinguish zero and non-zero storage slots
- What will happen if all storage slots with small values (less than a threshold  $U$ ) are stored separately from the entire tree? In this case,  $U$  reserved storage sets are enabled to hold all slots with value  $0, 1, \dots, U-1$
- Unsure if it actually improves performance for any  $U$

# Thanks for listening!

Questions welcome